

Yiting Qu — CV

✉ yiting.qu@cispa.de • [yitingqu.github.io](https://github.com/yitingqu)

Employment

CISPA Helmholtz Center for Information Security <i>Postdoctoral Researcher</i> <i>Hosted by Prof. Michael Backes and Prof. Yang Zhang</i> <i>Thesis: Mitigating Risks in Real-World and AI-Generated Visual Content</i>	Saarbrücken, Germany 11/2025 -
---	--

Education

CISPA Helmholtz Center for Information Security <i>Ph.D. in Computer Science (Awarded by Saarland University)</i> <i>Advisors: Prof. Michael Backes, Prof. Yang Zhang</i>	Saarbrücken, Germany 11/2021 - 10/2025
Shanghai Jiao Tong University <i>M.Sc. in Economics and Management</i> <i>Advisor: Prof. Suguo Du</i>	Shanghai, China 9/2018 - 6/2021
Shandong University <i>B.Sc. in Management</i> <i>Advisor: Prof. Tao Sun</i>	Jinan, China 9/2014 - 6/2018

Research Interests

- **Real-world misuse of generative AI:** unsafe image generation, harmful multimodal content, hateful memes, and AI-generated hate speech.
- **Security of generative models:** prompt stealing, jailbreaking, adversarial bypasses, and cross-modal safety gaps in vision-language models.
- **Mitigation and synthetic media forensics:** safety benchmarks, harmful-content detection, deepfake analysis, and fake point-cloud detection.

Publications

My full publication list is available on DBLP and Google Scholar (500+ citations as of May 2026). My ORCID is 0009-0000-1638-7287. IEEE S&P, CCS, USENIX Security, and NDSS are widely recognized as top-tier venues in information security.

Conference.....

- [1] **Yiting Qu**, Xinyue Shen, Yixin Wu, Michael Backes, Savvas Zannettou, and Yang Zhang. UnsafeBench: Benchmarking Image Safety Classifiers on Real-World and AI-Generated Images. In *ACM SIGSAC Conference on Computer and Communications Security (CCS)*. ACM, 2025 (**CCF-A**, **Acceptance rate: 14.5%**).
Integrated into Promptfoo, an AI safety testing platform under OpenAI.

- [2] **Yiting Qu**, Ziqing Yang, Yihan Ma, Michael Backes, and Yang Zhang. Hate in Plain Sight: On the Risks of Moderating AI-Generated Hateful Illusions. In *IEEE International Conference on Computer Vision (ICCV)*. ICCV, 2025 (CCF-A, Acceptance rate: 24.0%).
- [3] **Yiting Qu**, Michael Backes, and Yang Zhang. Bridging the Gap in Vision Language Models in Identifying Unsafe Concepts Across Modalities. In *USENIX Security Symposium (USENIX Security)*. USENIX, 2025 (CCF-A, Acceptance rate: 17.1%).
- [4] Xinyue Shen, Yixin Wu, **Yiting Qu**, Michael Backes, Savvas Zannettou, and Yang Zhang. HateBench: Benchmarking Hate Speech Detectors on LLM-Generated Content and Hate Campaigns. In *USENIX Security Symposium (USENIX Security)*. USENIX, 2025 (CCF-A, Acceptance rate: 17.1%).
Recognized by OpenAI's Promptfoo as a documented LLM vulnerability.
- [5] Yihan Ma, Xinyue Shen, **Yiting Qu**, Ning Yu, Michael Backes, Savvas Zannettou, and Yang Zhang. From Meme to Threat: On the Hateful Meme Understanding and Induced Hateful Content Generation in Open-Source Vision Language Models. In *USENIX Security Symposium (USENIX Security)*. USENIX, 2025 (CCF-A, Acceptance rate: 17.1%).
- [6] **Yiting Qu**, Zhikun Zhang, Yun Shen, Michael Backes, and Yang Zhang. FAKEPCD: Fake Point Cloud Detection via Source Attribution. In *ACM Asia Conference on Computer and Communications Security (ASIACCS)*. ACM, 2024 (CORE A, Acceptance rate: 22.0%).
- [7] Xinyue Shen, **Yiting Qu**, Michael Backes, and Yang Zhang. Prompt Stealing Attacks Against Text-to-Image Generation Models. In *USENIX Security Symposium (USENIX Security)*. USENIX, 2024 (CCF-A, Acceptance rate: 19.2%).
Recognized by Microsoft in the Microsoft Vulnerability Severity Classification; over 30K dataset downloads.
- [8] **Yiting Qu**, Xinyue Shen, Xinlei He, Michael Backes, Savvas Zannettou, and Yang Zhang. Unsafe Diffusion: On the Generation of Unsafe Images and Hateful Memes From Text-To-Image Models. In *ACM SIGSAC Conference on Computer and Communications Security (CCS)*. ACM, 2023 (CCF-A, Acceptance rate: 19.2%).
Cited in policy documents from the US, the UK, and Germany; Covered by Informationsdienst Wissenschaft (idw), The Next Web, and Montreal AI Ethics Institute.
- [9] **Yiting Qu**, Xinlei He, Shannon Pierson, Michael Backes, Yang Zhang, and Savvas Zannettou. On the Evolution of (Hateful) Memes by Means of Multimodal Contrastive Learning. In *IEEE Symposium on Security and Privacy (S&P)*. IEEE, 2023 (CCF-A, Acceptance rate: 17.4%).
- Journal.....
- [10] **Yiting Qu**, Suguo Du, Shaofeng Li, Yan Meng, Le Zhang, and Haojin Zhu. Automatic Permission Optimization Framework for Privacy Enhancement of Mobile Applications. *IEEE Internet of Things Journal*, 2020 (JCR Q1,IF=9.6).
- Technical Report.....
- [11] Mengfei Liang, **Yiting Qu**, Yukun Jiang, Michael Backes, and Yang Zhang. From Evidence to Verdict: An Agent-Based Forensic Framework for AI-Generated Image Detection. *CoRR abs/2511.00181*, 2025.

[12] Junjie Chu, **Yiting Qu**, Ye Leng, Michael Backes, Yun Shen, Savvas Zannettou, and Yang Zhang. Understanding LLM Behavior When Encountering User-Supplied Harmful Content in Harmless Tasks. *CoRR abs/2603.11914*, 2026.

Research Impact and Media Coverage

- 2026 May - UnsafeBench was integrated into *Promptfoo* as an official plugin.
- 2025 Oct - UK Department for Science, Innovation & Technology (DSIT). *International Scientific Report on the Safety of Advanced AI: Interim Report*
- 2025 Jan - German Federal Office for Information Security (BSI). *Generative AI Models: Opportunities and Risks for Industry and Authorities*.
- 2025 Jan - HateBench was incorporated into *Promptfoo* as an LLM hate campaign vulnerability.
- 2024 July - US. National Institute of Standards and Technology (NIST). *NIST Trustworthy and Responsible AI: NIST AI 600-1*
- 2024 July - Unsafe Diffusion was included in *Awesome-Diffusion-Models* with 11K+ GitHub stars.
- 2024 July - NetApp. *FAKEPCD: Fake Point Cloud Detection via Source Attribution*.
- 2024 Jan - Prompt stealing attack was recognized in the *Microsoft Vulnerability Severity Classification for AI Systems*.
- 2024 Jan - US. National Institute of Standards and Technology (NIST). *NIST Trustworthy and Responsible AI: NIST AI 100-2e2023*
- 2023 Oct - Informationsdienst Wissenschaft (idw). *AI Image Generators as Drivers of Unsafe Images? CISPA Researcher Develops Filter to Tackle This*.
- 2023 Jul - Montreal AI Ethics Institute (MAIEI). *On the Generation of Unsafe Images and Hateful Memes From Text-To-Image Models*.

Selected Honors and Awards

- 2024 - Excellence Scholarship for Overseas Self-financed PhD Students
- 2018 - National Scholarship at Shanghai Jiao Tong University
- 2016 - Weichai Scholarship at Shandong University

Teaching

- 2025, Co-Lecturer, Attacks Against Machine Learning Models, CISPA
- 2025, Teaching Assistant, Data-driven Understanding of the Disinformation Epidemic, CISPA
- 2024, Teaching Assistant, Attacks Against Machine Learning Models, CISPA
- 2024, Teaching Assistant, Data-driven Understanding of the Disinformation Epidemic, CISPA

- 2020, Teaching Assistant, Statistics and Machine Learning, Shanghai Jiao Tong University

Mentoring

- 2025-2026, Bo Shao, Ph.D. student, Computer Science, CISPA
- 2025-2026, Mengfei Liang, Ph.D. student, Computer Science, CISPA
- 2025-2026, Yibo Liang, Ph.D. student, Computer Science, Xi'an Jiaotong University
- 2025-2026, Yuheng Wang, B.S., Computer Science, Xi'an Jiaotong University
- 2024-2025 Junjie Chu, Ph.D. student, Computer Science, CISPA

Academic Service

Program Committees.....

- Privacy Enhancing Technologies Symposium (PETs/PoPETs), 2026, 2027
- Conference on Machine Learning and Systems (MLSys), 2026
- USENIX Security Symposium, Artifact Evaluation Committee (USENIX Security AEC), 2025

Journal Reviewing.....

- IEEE Transactions on Information Forensics and Security (TIFS), 2025, 2026
- IEEE Transactions on Dependable and Secure Computing (TDSC), 2025
- ACM Transactions on Privacy and Security (TOPS), 2024

Conference Reviewing.....

- ACM SIGSAC Conference on Computer and Communications Security (CCS), 2023, 2024, 2025
- ACM Web Conference (WWW), 2023, 2024, 2025
- ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD), 2023, 2024, 2026
- Annual Meeting of the Association for Computational Linguistics (ACL), 2025
- IEEE/CVF International Conference on Computer Vision (ICCV), 2025
- USENIX Security Symposium (USENIX Security), 2025
- IEEE Symposium on Security and Privacy (IEEE S&P), 2024
- European Conference on Computer Vision (ECCV), 2024
- IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2024
- International Conference on Learning Representations (ICLR), 2024